


Artificial Intelligence


Other Search and Optimization Methods

Russell and Norvig
Chapter 4

prediction
predict $f(x)$ for a given x




optimization
find an x so that $f(x)$ is maximal



is this search?
is this reinforcement learning?

playground example

gradient ascent (hillclimbing)



gradient ascent (hillclimbing)

```

function HILL-CLIMBING(problem) returns a solution state
inputs: problem, a problem
static: current, a node
           next, a node

current ← MAKE-NODE(INITIAL-STATE[problem])
loop do
  next ← a highest-valued successor of current
  if VALUE[next] < VALUE[current] then return current
  current ← next
end
    
```


$$x_i = x_i + \alpha \, d\text{VALUE}/dx_i$$

learning rate (small positive constant)

problems with hillclimbing


problems:

- local maxima -> random restarts
- plateaus -> random restarts
- ridges



problems with hillclimbing one solution: STAGE [Boyan and Moore]

1. remember the maxima of all hillclimbing applications to the given function



2. estimate a function of the maxima



3. Stage 1: use the ending point as a starting point for hillclimbing on the function of the maxima



4. Stage 2: use the ending point as a starting point for hillclimbing on the given function
this will be the starting point of hillclimbing in step 3 after the function of the maxima was re-estimated




5. go to 2

simulated annealing

annealing = the process of gradually cooling a liquid until it freezes
if the temperature is lowered sufficiently slowly,
the material will attain a lowest-energy (perfectly ordered) configuration

hillclimbing with going downhill (from time to time)



VLSI layout and lots of other applications

simulated annealing

VALUE[] = total energy of the atoms in the material
T = temperature

```

function SIMULATED-ANNEALING(problem, schedule) returns a solution state
inputs: problem, a problem
         schedule, a mapping from time to "temperature"
static: current, a node
         next, a node
         T, a "temperature" controlling the probability of downward steps

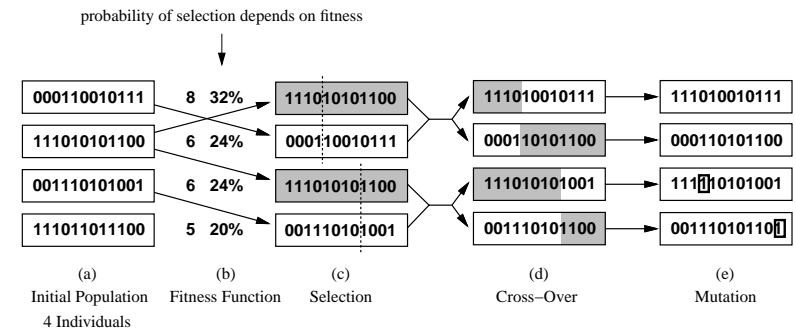
current ← MAKE-NODE(INITIAL-STATE(problem))
for t ← 1 to ∞ do
  T ← schedule[t]
  if T = 0 then return current
  next ← a randomly selected successor of current
  ΔE ← VALUE[next] − VALUE[current]
  if ΔE > 0 then current ← next
  else current ← next only with probability  $e^{\Delta E/T}$ 
    
```

genetic algorithms and evolutionary programming

hillclimbing with going downhill and parallel search

have a group of individuals
offsprings are “genetic mixtures” of their parents
some random mutations occur
if the performance of an individual is bad,
it dies early and thus doesn’t reproduce very often

genetic algorithms and evolutionary programming



? genetic algorithms are the third best way of doing just about anything ?