

Artificial Intelligence

Game Playing

Nilsson - Chapter 12
Russell and Norvig - Chapter 6



May 1997
Deep Blue - Garry Kasparov
3.5 - 2.5

machines are better than humans in:
othello

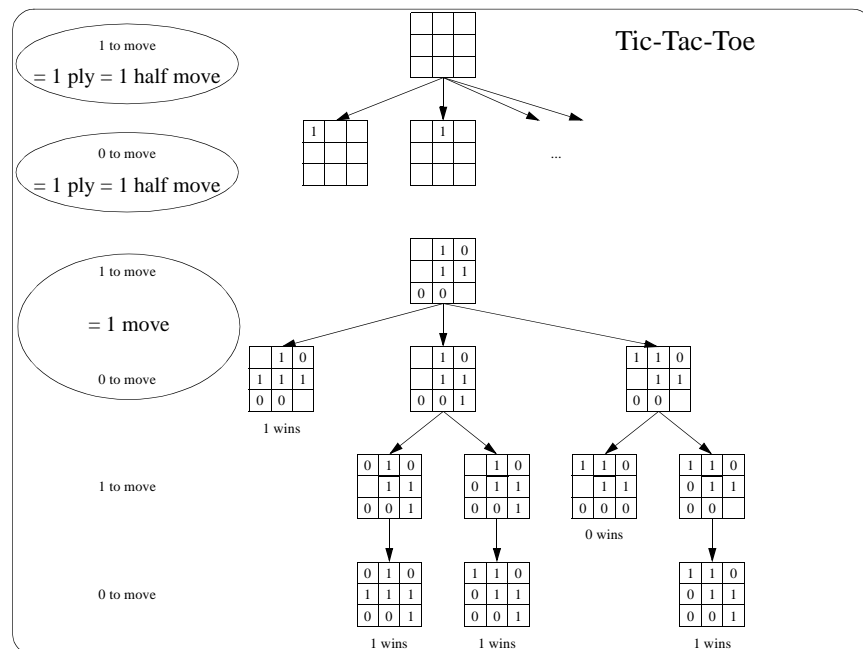
machines and humans are about equally good in:
checkers (draughts), backgammon, scrabble, chess, bridge?

humans are better than machines in:
go

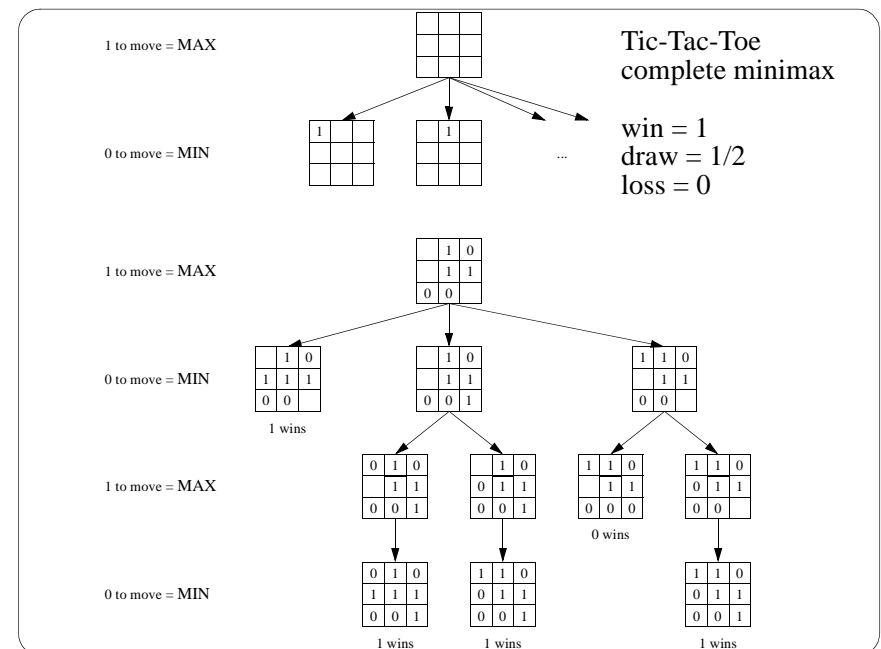
here: perfect information zero-sum games

Game Playing; page 1 of 15

Game Playing; page 2 of 15



Game Playing; page 3 of 15



Game Playing; page 4 of 15

minimax algorithm

at terminal node:
the minimax value of the terminal node is
given by the evaluation function (0 = loss, 1/2 = draw, 1 = win)

at a non-terminal MAX node:
calculate the minimax values of its successor nodes
the minimax value of the MAX node is
the maximum of the minimax values of its successor nodes

at a non-terminal MIN node:
calculate the minimax values of its successor nodes
the minimax value of the MIN node is
the minimum of the minimax values of its successor nodes

chess:
 10^{40} different legal positions
 35^{100} nodes in average game tree

issues

evaluation function

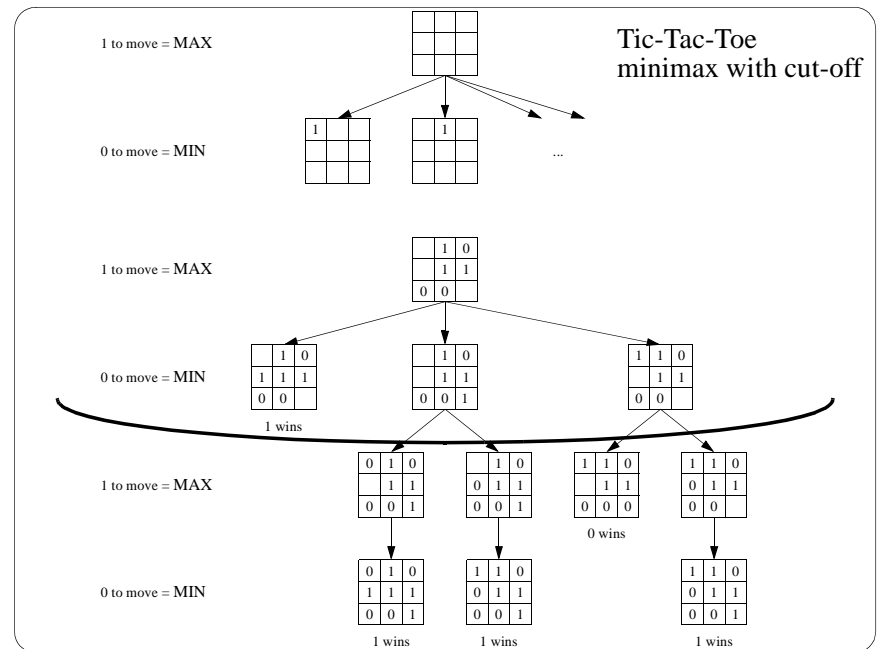
- roughly corresponds to the “likelihood of winning”
- must be between the value of a winning terminal state and a losing one
- must be easy to calculate (= fast)

often a weighted linear function
with handselected features and learned coefficients

features for chess

-
-
-
-
-

reach quiescence (horizon effect)



alpha-beta

= same result as minimax but more efficient

insight:


one does not need to look at all nodes to determine
the minimax value of the root of the game tree

alpha-beta algorithm

perform a depth-first minimax search


if the minimax value of the root of the game tree is the same
not matter what the minimax value of a node is
then we can prune this node and all of its descendants

an easy(?) example




often, alpha-beta is able to cut huge subtrees
ideally, it can search a game tree twice as deep as minimax
in the same amount of time

a harder(?) example



a harder(?) example: minimax

the minimax value of the
(root of the) game (tree) = 4



alpha-beta algorithm

alpha = best (largest) minimax value MAX is guaranteed to reach
beta = best (smallest) minimax value MIN is guaranteed to reach

call MAX-VALUE(state = root, alpha = loss, beta = win)

```


MAX-VALUE(state, alpha, beta)
if node is a terminal node (or to be treated like one) then
    return the value of the evaluation function for that node
else
    for each successor state s of state do
        set alpha := MAX(alpha, MIN-VALUE(s, alpha, beta))
        if alpha >= beta then return alpha
    return alpha
    
```

```


MIN-VALUE(state, alpha, beta)
if node is a terminal node (or to be treated like one) then
    return the value of the evaluation function for that node
else
    for each successor state s of state do
        set beta := MIN(beta, MAX-VALUE(s, alpha, beta))
        if alpha >= beta then return beta
    return beta
    
```

a harder(?) example: alpha-beta (a “deep” cut)


the minimax value of the
(root of the) game (tree) = 4



another example



another example



try “killer moves” early