


# Artificial Intelligence

## Markov Decision Problems

Nilsson - briefly mentioned in Chapter 10  
Russell and Norvig - Chapter 17

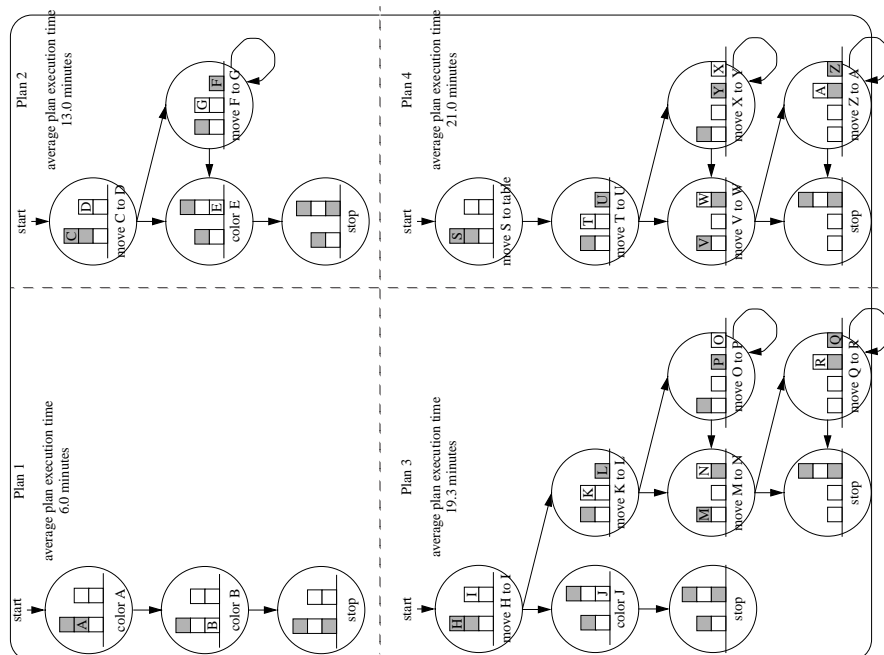
Markov Decision Problems; page 1 of 30

example: a probabilistic blockworld




action	outcome	probability	time (= cost)
move	success	0.1	1 minute
	failure	0.9	1 minute
paint	success	1.0	3 minutes

Markov Decision Problems; page 2 of 30



Markov Decision Problems; page 3 of 30

how to determine the average plan-execution time of a given plan



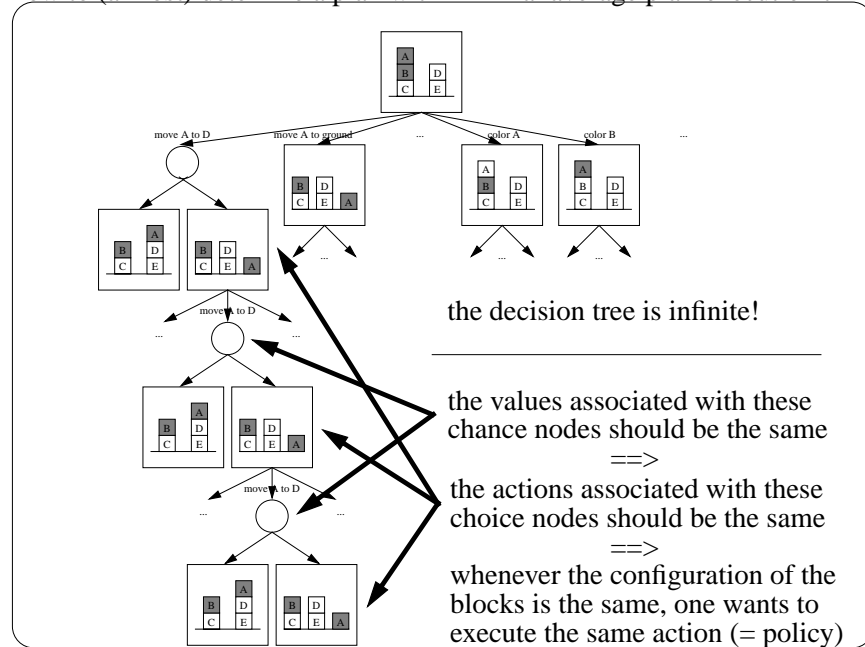
$t_i$  = average plan-execution time until a goal state is reached if the agent starts in state  $i$  and follows the plan

$$\begin{aligned} t_1 &= 1 + 0.1 t_2 + 0.9 t_3 \\ t_2 &= 3 + 1.0 t_4 \\ t_3 &= 1 + 0.1 t_2 + 0.9 t_3 \\ t_4 &= 0 \end{aligned}$$

$$\begin{aligned} t_1 &= 13 \quad (= \text{average plan-execution time}) \\ t_2 &= 3 \\ t_3 &= 13 \\ t_4 &= 0 \end{aligned}$$



Markov Decision Problems; page 4 of 30

how to (almost) determine a plan with minimal average plan-execution time



Markov Decision Problems; page 5 of 30


## deterministic planning and search



actions have deterministic effects  
state and action determine uniquely the successor state  
states are completely observable  
a plan is a sequence of actions (= path)  
minimize total cost  
optimal plan is acyclic

Markov Decision Problems; page 6 of 30

## probabilistic planning and search =Markov Decision Problems (MDPs)



Markov property

actions have probabilistic effects  
state and action uniquely determine prob distribution over successor states  
states are completely observable  
a plan is a mapping from states to actions (= policy)  
minimize expected total cost  
optimal plan can be cyclic

Markov Decision Problems; page 7 of 30

## deterministic planning and search



a plan is a sequence of actions (= path)  
can be found using (forward or backward) search

## Markov Decision Problems


a plan is a mapping from states to actions (= policy)  
how to find it?

determine the expected goal distances of all states

greedily assign the action to each state  
that decreases the expected goal distance the most

Markov Decision Problems; page 8 of 30

## deterministic planning and search



$s$  = state  
 $a$  = action  
 $A(s)$  = set of actions that can be executed in state  $s$   
 $\text{succ}(s,a)$  = the state that results from the execution of action  $a$  in state  $s$   
 $c(s,a)$  = the cost that results from the execution of action  $a$  in state  $s$


$\text{gd}(s)$  = goal distance of state  $s$

$\text{gd}(s) = 0$  if  $s$  is a goal state  
 $\text{gd}(s) = \min_{a \in A(s)} c(s,a) + \text{gd}(\text{succ}(s,a))$  if  $s$  is not a goal state

$a(s)$  = the optimal action to execute in state  $s$

$a(s) = \text{argmin}_{a \in A(s)} c(s,a) + \text{gd}(\text{succ}(s,a))$  if  $s$  is not a goal state

## Markov Decision Problems



$s$  = state  
 $a$  = action  
 $A(s)$  = set of actions that can be executed in state  $s$   
 $\text{succ}(s,a)$  = the set of states that can result from the execution of action  $a$  in state  $s$   
 $c(s,a)$  = the cost that results from the execution of action  $a$  in state  $s$   
 $p(s'|s,a)$  = the probability that state  $s'$  results from the execution of action  $a$  in state  $s$   
 $\text{gd}(s)$  = expected goal distance of state  $s$

### Bellman equation


$\text{gd}(s) = 0$  if  $s$  is a goal state  
 $\text{gd}(s) = \min_{a \in A(s)} (c(s,a) + \sum_{s' \in \text{succ}(s,a)} p(s'|s,a) \text{gd}(s'))$  if  $s$  is not a goal state

$a(s)$  = the optimal action to execute in state  $s$


$a(s) = \text{argmin}_{a \in A(s)} (c(s,a) + \sum_{s' \in \text{succ}(s,a)} p(s'|s,a) \text{gd}(s'))$  if  $s$  is not a goal state

## examples

### deterministic planning and search




### Markov Decision Problems



given the expected goal distances, we can use the definition to check them but calculating them is a chicken-and-egg problem

## deterministic planning and search



$s$  = state  
 $a$  = action  
 $A(s)$  = set of actions that can be executed in state  $s$   
 $\text{succ}(s,a)$  = the state that results from the execution of action  $a$  in state  $s$   
 $c(s,a)$  = the cost that results from the execution of action  $a$  in state  $s$


$\text{gd}(s)$  = goal distance of state  $s$  (= minimal cost until a goal is reached if execution starts in state  $s$ )  
 $\text{gd}_i(s)$  = minimal cost until a goal is reached or  $i$  actions have been executed if execution starts in state  $s$

for  $i$  larger than a constant:  $\text{gd}(s) = \text{gd}_i(s)$  (= once  $\text{gd}_i(s) = \text{gd}_{i-1}(s)$  for all states  $s$ )

$\text{gd}_0(s) = 0$

$\text{gd}_i(s) = 0$  if  $s$  is a goal state  
 $\text{gd}_i(s) = \min_{a \in A(s)} (c(s,a) + \text{gd}_{i-1}(\text{succ}(s,a)))$  if  $s$  is not a goal state

## Markov Decision Problems



$s$  = state  
 $a$  = action  
 $A(s)$  = set of actions that can be executed in state  $s$   
 $\text{succ}(s,a)$  = the set of states that can result from the execution of action  $a$  in state  $s$   
 $c(s,a)$  = the cost that results from the execution of action  $a$  in state  $s$   
 $p(s'|s,a)$  = the probability that state  $s'$  results from the execution of action  $a$  in state  $s$   
 $gd(s)$  = expected goal distance of state  $s$   
 $gd_i(s)$  = minimal expected cost until a goal is reached or  $i$  actions have been executed if execution starts in state  $s$

$gd(s) = \lim_{i \rightarrow \infty} gd_i(s)$  (not necessarily after a finite amount of time)

$gd_0(s) = 0$

$gd_i(s) = 0$  if  $s$  is a goal state  
 $gd_i(s) = \min_{a \in A(s)} (c(s,a) + \sum_{s' \in \text{succ}(s,a)} p(s'|s,a) gd_{i-1}(s'))$  if  $s$  is not a goal state

Markov Decision Problems; page 13 of 30

## Value Iteration

maintains approximations of the goal distances (= values)

1.  $i := 0$
2. Set (for all  $s \in S$ )  $gd_i(s) = 0$ .
3.  $i := i+1$
4. Set (for all  $s \in S$ )
 


$gd_i(s) = 0$   
 $gd_i(s) = \min_{a \in A(s)} (c(s,a) + \sum_{s' \in \text{succ}(s,a)} p(s'|s,a) gd_{i-1}(s'))$

if  $s$  is a goal state  
 if  $s$  is not a goal state
5. If (for some  $s \in S$ )  $|gd_i(s) - gd_{i-1}(s)| > \text{small constant}$ , go to 3.
6. Set (for all  $s \in S$  that are not goal states)
 

$a(s) = \text{argmin}_{a \in A(s)} (c(s,a) + \sum_{s' \in \text{succ}(s,a)} p(s'|s,a) gd_i(s'))$

Markov Decision Problems; page 14 of 30


## example of Value Iteration (1)



i		0	1	2	3	4	5	6			
start	a1	0	3	1	3	2	3	2.75	3	3	3
	a2	0	3	1	2	2	2.75	2.75	3	3.75	3
state 1	a3	0	1	1	1.5	1.5	2	2	2.375	2.375	2.5
		0	1	1	1.5	1.5	2	2	2.375	2.375	2.5
goal		0	0	0	0	0	0	0	0	0	0

Markov Decision Problems; page 15 of 30

## example of Value Iteration (2)



(no discounting)

which action to execute in the start state?

$a1: 3 + 0 = 3$   
 $a2: 1 + 0.5 \cdot 3 + 0.5 \cdot 2.5 = 3.75$

execute a1!

Markov Decision Problems; page 16 of 30


## Policy Iteration

maintains a policy

1.  $i := 0$
2. Set (for all  $s \in S$  that are not goal states)  $a_i(s)$  to an arbitrary action in  $A(s)$ .
3. Set  $gd_i(s)$  to the average plan-execution time until a goal state is reached if the agent starts in state  $s$  and follows policy  $a_i$
4.  $i := i+1$
5. Set (for all  $s \in S$  that are not goal states)
 
$$a_i(s) = \operatorname{argmin}_{a \in A(s)} (c(s,a) + \sum_{s' \in \operatorname{succ}(s,a)} p(s'|s,a) gd_i(s'))$$
6. If (for some  $s \in S$  that is not a goal state)  $a_i(s)$  does not equal  $a_{i-1}(s)$ , go to 3
7. Set (for all  $s \in S$  that are not goal states)  $a(s) = a_i(s)$ .

Note: The initial policy  $a_0$  has to guarantee that the agent reaches a goal state with probability one no matter which state it is started in.

## example of Policy Iteration



policy at  $i=0$   
 $a_0(\text{start}) = a_2$  (could also have been  $a_1$ ) and  $a_0(\text{state } 1) = a_3$

$$\begin{aligned} gd_0(\text{start}) &= 1 + 0.5 gd_0(\text{start}) + 0.5 gd_0(\text{state } 1) = 6 \\ gd_0(\text{state } 1) &= 1 + 0.5 gd_0(\text{start}) + 0.5 gd_0(\text{goal}) = 4 \\ gd_0(\text{goal}) &= 0 \end{aligned}$$

policy at  $i=1$   
 $a_1(\text{start}) = a_1$  and  $a_1(\text{state } 1) = a_3$

$$\begin{aligned} gd_1(\text{start}) &= 3 + 1.0 gd_0(\text{goal}) = 3 \\ gd_1(\text{state } 1) &= 1 + 0.5 gd_1(\text{start}) + 0.5 gd_1(\text{goal}) = 2.5 \\ gd_1(\text{goal}) &= 0 \end{aligned}$$

policy at  $i=2$   
 $a_2(\text{start}) = a_1$  and  $a_2(\text{state } 1) = a_3$

execute action  $a_1$  in the start state!

## extensions: no goal (1)

what if there is no goal?  
 “living in the world”

can no longer minimize expected cost until the goal is reached

- here:
- can minimize expected cost per action execution
  - can minimize expected total discounted cost

## extensions: no goal (2)

cannot minimize expected total cost

1	2	3	4	4	4	4	...	expected total cost = infinite
1	1	1	1	1	1	1	...	expected total cost = infinite

extensions: no goal (3)

discount factor

total discounted cost =



if the interest rate is  $(1-\gamma)/\gamma$  (for  $0 < \gamma < 1$ ),  
how much money do I need to pay someone right now  
so that there is no difference to paying the following yearly installments

1   2   3   4   4   4   4   ...

x dollars right now are worth  $(1 + (1-\gamma)/\gamma)x = x/\gamma$  dollars in a year  
so, y dollars in a year are worth  $\gamma y$  dollars right now

answer:  $1 + \gamma 2 + \gamma^2 3 + \gamma^3 4 + \gamma^4 4 + \dots$

extensions: no goal (4)

can minimize the expected total discounted cost - assume  $\gamma = 0.9$

1   2   3   4   4   ...   expected total discounted cost = 34.39  
1   1   1   1   1   ...   expected total discounted cost = 10.00

extensions: no goal (5)

- discounting makes the total cost finite

c   c   c   c   c   ...   expected total discounted cost =  $c/(1-\gamma)$

- discounting smoothes out the horizon

- discounting can be interpreted as the probability of dying

discounting:

if the interest rate is  $(1-\gamma)/\gamma$ ,  
then y dollars in a year are worth  $\gamma y$  dollars right now


dying:

if I die later this year with probability  $1-\gamma$ ,  
then the expected value of y dollars in a year is  $\gamma y$  right now

### Value-Iteration with or without discounting

$\gamma$  = discount factor ( $0 < \gamma < 1$ ); if there is a goal, can set  $\gamma = 1$  (no discounting)  
 $s$  = state  
 $a$  = action  
 $A(s)$  = set of actions that can be executed in state  $s$   
 $\text{succ}(s,a)$  = the set of states that can result from the execution of action  $a$  in state  $s$   
 $c(s,a)$  = the cost that results from the execution of action  $a$  in state  $s$   
 $p(s'|s,a)$  = the probability that state  $s'$  results from the execution of action  $a$  in state  $s$   
 $gd(s)$  = minimal expected discounted total cost if execution starts in state  $s$   
 $gd(s) = 0$  if  $s$  is a goal state  
 $gd(s) = \min_{a \in A(s)} (c(s,a) + \gamma \sum_{s' \in \text{succ}(s,a)} p(s'|s,a) gd(s'))$  if  $s$  is not a goal state  
 $gd_i(s)$  = minimal expected discounted total cost until a goal is reached or  $i$  actions have been executed if execution starts in state  $s$   
 $gd_0(s) = 0$  for all  $s$   
 $gd_i(s) = 0$  if  $s$  is a goal state  
 $gd_i(s) = \min_{a \in A(s)} (c(s,a) + \gamma \sum_{s' \in \text{succ}(s,a)} p(s'|s,a) gd_{i-1}(s'))$  if  $s$  is not a goal state  
 $gd(s) = \lim_{i \rightarrow \infty} gd_i(s)$   
 $a(s)$  = the optimal action to execute in state  $s$   
 $a(s) = \text{argmin}_{a \in A(s)} (c(s,a) + \gamma \sum_{s' \in \text{succ}(s,a)} p(s'|s,a) gd(s'))$  if  $s$  is not a goal state  
 $gd(s)$  does not necessarily converge after a finite amount of time  
 $a(s)$  converges after a finite amount of time if  $gd(s)$  is approximated with  $gd_i(s)$  for all  $s$

### example of Value Iteration (1)




i		0	1	2	3	4
start	a1	0	3	1	3	1.9
	a2	0	1	1	3	1.9
state 1	a3	0	1	1	1.45	1.45
	a1	0	0	0	0	0
goal	a1	0	0	0	0	0
	a2	0	0	0	0	0

	5	6	7
start	3	3	3
state 1	3.2912	3.4	3.4075
goal	2.3334	2.3334	2.35
	0	0	0

Markov Decision Problems; page 25 of 30

### example of Value Iteration (2)



which action to execute in the start state?

$$a1: 3 + 0 = 3$$

$$a2: 1 + 0.9 \cdot 0.5 \cdot 3 + 0.9 \cdot 0.5 \cdot 2.35 = 3.4075$$


execute a1!

(In general, the optimal action depends on the discount factor!)

Markov Decision Problems; page 26 of 30

### “learning” for optimization

“reinforcement learning” with Markov Decision Process Models



exam example

find a policy (behavior)  
that maximizes the expected total discounted reward  
even in the presence of delayed rewards

if you don't know the action outcomes (rewards and probabilities):  
reinforcement learning

exploration/exploitation tradeoff


Markov Decision Problems; page 27 of 30

### “learning” for optimization

“reinforcement learning” with Markov Decision Process Models

approach 1

estimate the probabilities and rewards



use value-iteration

Markov Decision Problems; page 28 of 30

“learning” for optimization

“reinforcement learning” with Markov Decision Process Models

approach 2

use Q-learning

if you execute action  $a$  in state  $s$  and  
you receive cost  $c$  and make a transition to state  $s'$   
then update

$$Q(s,a) = Q(s,a) + \alpha (c + \gamma V(s') - Q(s,a))$$

$$V(s') = \min_{a \in A(s')} Q(s',a)$$

$\uparrow$  learning rate  
 $\uparrow$  discount factor  
 $0 < \gamma < 1$

$Q(s,a)$  = minimal expected discounted total cost until a goal is reached  
if execution starts in state  $s$  and the first action executed is  $a$   
 $V(s')$  = minimal expected discounted total cost until a goal is reached  
if execution starts in state  $s'$  (= “ $gd(s')$ ”)

“learning” for optimization

“reinforcement learning” with Markov Decision Process Models

approach 2

1. Initialize  $Q(s,a) = 0$  for all states  $s$  and actions  $a$ .
2.  $s :=$  the current state.
3. if  $s$  is a goal state then stop.
4. Choose an action  $a$  to execute in the current state  $s$ .  
(The action believed to be best is  $a := \operatorname{argmin}_{a \in A(s)} Q(s,a)$ .)
5. Execute action  $a$ . Observe the cost  $c$  and successor state  $s'$ .
6. Update  $Q(s,a) = Q(s,a) + \alpha (c + \gamma V(s') - Q(s,a))$ .
7. Goto 2.

